

cluster-ad0a - cluster-ad0a - Dat x +

console.cloud.google.com/dataproc/clusters/cluster-ad0a/instances?region=us-east4&authuser=1&folder=&project=wor...

Google Cloud Platform WordCount Search products and resources

Dataproc cluster-ad0a SUBMIT JOB REFRESH DELETE VIEW LOGS Google account: Brendan Russ (bvr6@pitt.edu)

Clusters Jobs Workflows Autoscaling policies Component exchange Notebooks

For PD-Standard without local SSDs, we strongly recommend provisioning 1TB or larger to ensure consistently high I/O performance. See <https://cloud.google.com/compute/docs/disks/performance> for information on disk I/O performance.

Name	cluster-ad0a
Cluster UUID	d5baf40-2bb5-446d-b0ca-551ac079dc44
Type	Dataproc Cluster
Status	Running

MONITORING JOBS VM INSTANCES CONFIGURATION WEB INTERFACES

Filter instances

Name	Role
cluster-ad0a-m	Master
cluster-ad0a-w-0	Worker
cluster-ad0a-w-1	Worker

Equivalent REST

```
bvr6@cluster-ad0a-m: ~/CloudLeas5 - Google Chrome
ssh.cloud.google.com/projects/wordcount-293603/zones/us-east4-c/instances/cluster-ad0a-m?authuser=1&hl=en_US&projectNumber=109974008917&useAdminProxy=true
Connected, host fingerprint: ssh-rsa 0 D6:AF:1A:F5:0C:29:01:79:9F:47:D0:61:6E:139
:B2:14:0A:85:1D:A3:9D:C3:38:27:E6:2D:7E:63:6E:59:25:AD
Linux cluster-ad0a-m 5.8.0-0.bpo.2-amd64 #1 SMP Debian 5.8.10-1-bpo10+1 (2020-09
-26) x86_64

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/*copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.
Last login: Tue Nov 10 01:43:03 2020 from 35.235.241.32
bvr6@cluster-ad0a-m:~$ ls
CloudLeas5 CloudWordCount
bvr6@cluster-ad0a-m:~$ cd CloudLeas5
bvr6@cluster-ad0a-m:~/CloudLeas5$ ls
CloudLeas5.jar bin cloudLeas5 lib src wordData
bvr6@cluster-ad0a-m:~/CloudLeas5$ hadoop jar CloudLeas5.jar wordData wordOutput
20/11/10 03:07:18 INFO client.RMProxy: Connecting to ResourceManager at cluster-ad0
a-m/10.150.0.4:8032
20/11/10 03:07:19 INFO client.AHSProxy: Connecting to Application History server at
cluster-ad0a-m/10.150.0.4:10200
20/11/10 03:07:19 INFO input.FileInputFormat: Total input files to process : 3
20/11/10 03:07:19 WARN hdfs.DataStreamer: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at org.apache.hadoop.hdfs.DataStreamer.closeResponder(DataStreamer.java:980
)
    at org.apache.hadoop.hdfs.DataStreamer.endBlock(DataStreamer.java:630)
    at org.apache.hadoop.hdfs.DataStreamer.run(DataStreamer.java:807)
20/11/10 03:07:19 INFO mapreduce.JobSubmitter: number of splits:3
20/11/10 03:07:20 INFO Configuration.deprecation: yarn.resourcemanager.system-metri
cs-publisher.enabled is deprecated. Instead, use yarn.system-metrics-publisher.enab
led
20/11/10 03:07:20 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_16036
77175143_0033
20/11/10 03:07:20 INFO impl.YarnClientImpl: Submitted application application_16036
77175143_0033
20/11/10 03:07:20 INFO mapreduce.Job: The url to track the job: http://cluster-ad0a
-m:8088/proxy/application_1603677175143_0033/
```

```
bvr6@cluster-ad0a-m: ~/CloudLeas5 - Google Chrome
ssh.cloud.google.com/projects/wordcount-293603/zones/us-east4-c/instances/cluster-ad0a-m?authuser=1&hl=en_US&projectNumber=109974008917&useAdminProxy=true

20/11/10 03:07:20 INFO mapreduce.Job: Running job: job_1603677175143_0033
20/11/10 03:07:27 INFO mapreduce.Job: Job job_1603677175143_0033 running in uber mo
de : false
20/11/10 03:07:27 INFO mapreduce.Job: map 0% reduce 0%
20/11/10 03:07:40 INFO mapreduce.Job: map 100% reduce 0%
20/11/10 03:07:48 INFO mapreduce.Job: map 100% reduce 20%
20/11/10 03:07:49 INFO mapreduce.Job: map 100% reduce 60%
20/11/10 03:07:51 INFO mapreduce.Job: map 100% reduce 80%
20/11/10 03:07:52 INFO mapreduce.Job: map 100% reduce 100%
20/11/10 03:07:52 INFO mapreduce.Job: Job job_1603677175143_0033 completed successfully
20/11/10 03:07:52 INFO mapreduce.Job: Counters: 50

File System Counters
  FILE: Number of bytes read=1410
  FILE: Number of bytes written=1666421
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=1027
  HDFS: Number of bytes written=156
  HDFS: Number of read operations=34
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=15

Job Counters
  Killed reduce tasks=1
  Launched map tasks=3
  Launched reduce tasks=5
  Data-local map tasks=3
  Total time spent by all maps in occupied slots (ms)=116224
  Total time spent by all reduces in occupied slots (ms)=143328
  Total time spent by all map tasks (ms)=29056
  Total time spent by all reduce tasks (ms)=35832
  Total vcore-milliseconds taken by all map tasks=29056
  Total vcore-milliseconds taken by all reduce tasks=35832
  Total megabyte-milliseconds taken by all map tasks=59506688
  Total megabyte-milliseconds taken by all reduce tasks=73383936

Map-Reduce Framework
  Map input records=3
  Map output records=116
  Map output bytes=1148
  Map output materialized bytes=1470
  Input split bytes=345
  Combine input records=0

bvr6@cluster-ad0a-m:~/CloudLeas5$
```

```
bvr6@cluster-ad0a-m: ~/CloudLeas5 - Google Chrome
ssh.cloud.google.com/projects/wordcount-293603/zones/us-east4-c/instances/cluster-ad0a-m?authuser=1&hl=en_US&projectNumber=109974008917&useAdminProxy=true

Data-local map tasks=3
Total time spent by all maps in occupied slots (ms)=116224
Total time spent by all reduces in occupied slots (ms)=143328
Total time spent by all map tasks (ms)=29056
Total time spent by all reduce tasks (ms)=35832
Total vcore-milliseconds taken by all map tasks=29056
Total vcore-milliseconds taken by all reduce tasks=35832
Total megabyte-milliseconds taken by all map tasks=59506688
Total megabyte-milliseconds taken by all reduce tasks=73383936

Map-Reduce Framework
  Map input records=3
  Map output records=116
  Map output bytes=1148
  Map output materialized bytes=1470
  Input split bytes=345
  Combine input records=0
  Combine output records=0
  Reduce input groups=58
  Reduce shuffle bytes=1470
  Reduce input records=116
  Reduce output records=19
  Spilled Records=232
  Shuffled Maps =15
  Failed Shuffles=0
  Merged Map outputs=15
  GC time elapsed (ms)=1740
  CPU time spent (ms)=5260
  Physical memory (bytes) snapshot=2534617088
  Virtual memory (bytes) snapshot=27758739456
  Total committed heap usage (bytes)=2008547328

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=682
File Output Format Counters
  Bytes Written=156

bvr6@cluster-ad0a-m:~/CloudLeas5$
```

```
bvr6@cluster-ad0a-m:~/CloudLeast5$ hadoop fs -getmerge wordOutput collectedResults
bvr6@cluster-ad0a-m:~/CloudLeast5$ cat collectedResults
yep      18
her      3
security      2
with     1
soup     3
renew    2
told     1
the      4
client   3
event    2
was      1
simple    4
suspend  3
should   2
world    1
enough   5
government      3
of        2
threaten      1
bvr6@cluster-ad0a-m:~/CloudLeast5$
```